



• CONFIDENTIAL — FOR NESTLÉ INTERNAL USE

SnapLogic for Nestlé

Use Case Assessment & Platform Overview

A use case assessment prepared by SnapLogic, to be validated jointly by Sriven and Nestlé — combining SnapLogic platform expertise with deep SAP environment knowledge.

Date	Prepared by	Attendees
March 2026	SnapLogic & Sriven	Nestlé, Sriven & SnapLogic

Nestlé Contacts

Renan Costa, Venkat Alikatte, Naresh Chavva, Lucas Arroyo

8

USE CASES ASSESSED

100MB/s

PROVEN THROUGHPUT

7+

CUSTOMER REFERENCES

75%

FASTER INTEGRATION

Table of Contents

1 Executive Summary	2 Strategic Alignment
3 Use Case Assessments (8)	4 Platform Architecture & Capabilities
5 Customer Reference Portfolio	6 Summary
A Appendix A — Platform Architecture & Glossary	

Executive Summary

Nestlé is at an integration inflection point. With EMEA S/4HANA migration complete and Europe/AOA going live in April 2026, a strategic commitment to Snowflake as the analytics platform, and three explicit 2026 priorities stated on March 12 — **Augmented Integration** (AI-assisted asset production), **Unstructured Data Strategy** (SharePoint, IoT, PDFs, videos with no current central strategy), and **New Velocities** (real-time SAP-Snowflake data with fewer delays) — the right platform choice will define the speed and cost of Nestlé's digital operations for the next decade.

This document outlines SnapLogic's approach to each of the 8 identified use cases, grounded in the Sigma Framework and verified customer implementations. The delivery model is a deliberate partnership: **Striven** brings certified SnapLogic implementation expertise; **Nestlé** brings the SAP environment knowledge that no external party can replicate. Together they form the team needed to take these use cases from assessment to working integration — each side contributing where they are strongest.

The proposed next step is to provision a sandbox instance where Striven and Nestlé will jointly validate the proposed approach against Nestlé's actual environment. UC #2 (SAP APO → Snowflake bulk replication) is the confirmed Priority #1 starting point — the blueprint for all subsequent data integration flows.

Scope Caveat — From SnapLogic: Based on our understanding of your environment and the SAP solutions you are using, we expect to be able to support the use cases as detailed in this document. By testing our proposed solution to each use case against your sandbox environments, Striven and Nestlé will validate the approach and confirm fit before any production commitment.

4-6 days

NESTLÉ CURRENT DEV
TIME

2 hrs

SNAPLOGIC (MIKRON
REFERENCE)

1000+

PRE-BUILT
CONNECTORS (SNAPS)

100 MB/s

PROVEN SAP CDC
THROUGHPUT

Strategic Alignment with Nestlé Priorities

Based on Nestlé's stated priorities, here is how the SnapLogic platform — implemented by Sriven in collaboration with Nestlé's SAP team — directly addresses each strategic theme.

Augmented Integration & Reusability

SnapLogic's AI-assisted pipeline generation (SnapGPT), template library, and reusable Shared Components directly reduce the time to build new integrations. Citizen developer capabilities allow business teams to self-serve without replicating logic across environments.

AI to Lower Integration Costs

SnapGPT copilot generates pipeline configurations from natural language. AutoSync provides no-code data synchronization for standard patterns. Combined with the 1000+ pre-built Snaps, this reduces the specialized engineering headcount required for integration projects.

Unstructured Data Strategy (2026 Priority)

Nestlé confirmed on March 12 that unstructured data strategy is a top 2026 priority — images, videos, PDFs, SharePoint documents, IoT logs from SmartFactory devices — with no central strategy today. SnapLogic provides SharePoint Online Snap Pack, OPCUA/MQTT broker connectivity, metadata extraction pipelines, and SnapGPT-assisted pipeline generation to establish a reusable, governed unstructured data foundation that teams can consume rather than building individually.

Real-Time Data Access (SAP ↔ Snowflake)

Ultra Pipelines with Kafka Consumer/Producer Snaps deliver sub-second latency for real-time SAP-to-Snowflake replication. Snowpipe Streaming Snap enables continuous, low-latency ingestion. Proven at 100 MB/s sustained throughput at Siemens scale.

Governance & Reliability

Role-Based Access Control (RBAC), 90-day audit logging, pipeline validation, approval workflows, and centralized credential management (HashiCorp Vault, AWS Secrets Manager) provide enterprise-grade governance without slowing development velocity.

Lowest TCO vs. ADF / SNP Kyano

SnapLogic consolidates multiple tools (ADF, BODS, SNP Kyano Glue, custom scripts) into one platform. This eliminates multi-vendor licensing, reduces operational overhead, and cuts integration development time by 75% based on comparable FMCG customer implementations. Also addresses the ELT/ETL cost tension: SnapLogic's on-prem Groundplex can handle pre-Snowflake transformations where appropriate, reducing unnecessary Snowflake compute spend.

MuleSoft Coexistence — Data Integration, Not Replacement

Nestlé already has MuleSoft for API management and transactional integrations — and that is not the focus here. SnapLogic is additive: filling the data integration gap (bulk replication, CDC, streaming, unstructured data) where MuleSoft is not the right tool. SnapLogic pipelines can consume MuleSoft-published APIs where appropriate. These platforms serve different layers of the integration stack.

Use Case Assessments

Detailed SnapLogic approach for each of the 8 identified use cases, with validated implementation details and customer references.

1

SAP APO → OMP (Azure Data Lake)

CURRENT STATE

Source: SAP APO via SAP Data Services (BODS)

Target: Azure Data Lake → OMP

Frequency: Daily batch

Complexity: Multiple ABAP programs/functions for value calculations

SNAPLOGIC APPROACH

Replace BODS downstream processing with SnapLogic Groundplex on-premise. **Sriven** will configure SnapLogic pipelines to connect to the extraction interface defined by Nestlé's SAP/APO team (custom BAPI, BW extractor, or CIF interface), handling transformation, routing to Azure Data Lake, and scheduling. Complex field mapping and value calculations run through Mapper and Script Snaps in the pipeline. **Note:** defining the correct SAP APO extraction method is Nestlé's ABAP team's responsibility — Sriven and Nestlé align on this together during sandbox validation. This is true regardless of which integration platform is used and is not a SnapLogic-specific constraint.

KEY SNAPLOGIC COMPONENTS

SAP RFC Snap

SAP BAPI Execute

Mapper Snap

Script Snap (JS)

Azure Data Lake Write

Groundplex (On-Prem)

VALUE DELIVERED

- Eliminates SAP BODS licensing and maintenance for the post-extraction transformation layer
- Sriven implements visual pipeline design to replace BODS transformation logic — ABAP extraction itself remains with Nestlé's SAP team (this is true regardless of integration platform)
- Daily batch scheduling via SnapLogic Task Scheduler — replaces BODS job scheduling
- Full audit logging of transformation steps: what was processed, when, and what was loaded

Sigma Framework Best Practice: For complex ABAP-sourced transformations, use modular pipeline design with SnapLogic's Pipeline Execute Snap to separate extraction, transformation, and load phases. This enables independent testing and reuse across similar SAP source patterns — directly addressing Nestlé's reusability goal. The modular approach also means that if the extraction interface changes (e.g., BAPI updated), only the extraction pipeline needs updating, not the full transformation chain.

BLUEPRINT USE CASE

CURRENT STATE

Source: SAP APO custom table (hundreds of millions of rows)

Current Tool: Azure Data Factory (ADF)

Target: Snowflake STG Layer

Challenge: Partitioned queries needed, requires on-prem runtime, critical performance

SNAPLOGIC APPROACH

SnapLogic Groundplex deployed on-prem replaces ADF for SAP APO table extraction. For APO data resident in standard DB-layer Z-tables, the SAP RFC Snap (BAPI/custom function module) handles high-volume partitioned reads — configurable partition parameters replace ADF's custom parameterization logic. For liveCache-resident APO planning objects, extraction interfaces must be confirmed by Nestlé's SAP/APO team first (see architecture note below); SnapLogic handles everything downstream. Snowflake Bulk Load Snap delivers to staging with native Snowpipe Streaming available for lower-latency options. Built-in parallelization eliminates the custom ADF pipeline complexity.

KEY SNAPLOGIC COMPONENTS

SAP RFC Snap (BAPI/FM)

SAP S/4HANA Batch Read

Snowflake Bulk Load

Snowflake Snowpipe Streaming


Groundplex (On-Prem)

Partition Handler

BOTTLENECK RESOLUTION

- Built-in partition handling replaces custom ADF parameterization
- Groundplex on-prem meets network security and data residency requirements
- Performance reconciliation via SnapLogic monitoring dashboards
- Parameterised pipeline architecture enables load resumption and reuse across table replications — directly addressing Venkat's confirmed requirement

REFERENCE ARCHITECTURE: HIGH-VOLUME SAP → SNOWFLAKE



Customer Reference — Mikron (Food Manufacturing): Replaced ADF-based SAP-to-Snowflake pipeline. Integration time reduced from 2-3 days to 2 hours (75% reduction). Team size reduced from 3-5 people to 1-2. Native SAP-Snowflake connectivity on 3 medium nodes (6 cores). **This is the exact architecture Nestlé needs.**

Also reference — Nielsen (CPG): Met aggressive Q1 SAP go-live deadline without external consultants. Internal teams now manage all SAP integrations — directly addressing Nestlé's goal of reducing development time from 4-6 days.

March 12 Session — Explicitly Confirmed Requirements for UC #2:

Venkat (March 12): *"The load identification and parameterization for reuse — we don't want to rebuild again. We should be able to pre-trigger or reuse that setup for triggering another load."*

This is a direct requirement for SnapLogic's **parameterised pipeline architecture**: each pipeline run logs which partition/load was processed (using Pipeline Properties or a Snowflake metadata table), enabling subsequent runs to resume intelligently. Partition parameters (e.g., date range, key range) are externalised as pipeline parameters — not hardcoded — so the same pipeline template handles every table in the batch without rebuild.

Multi-million row scale confirmed. Naresh/Venkat described table-level replication from SAP via SLT (DB triggers → logging table → target), processing individual records through to bulk loads. SnapLogic's Groundplex JDBC connector can read SLT's logging tables directly as an alternative extraction mechanism — see UC #5 for the full SLT architecture detail.

SAP APO Data Extraction — Important Architecture Note

SAP APO is architecturally different from SAP ECC and S/4HANA. APO planning data (demand planning, supply network planning, production planning/DS) is frequently stored in **liveCache** — an in-memory object store that is not a relational database table. This means standard RFC_READ_TABLE *does not apply* to most APO planning objects.

The correct extraction approach depends on what Nestlé's APO team confirms about their specific tables:

- **liveCache objects** (most APO planning data) — must be extracted via dedicated APO BAPIs (e.g., BAPI_APO_*) or CIF outbound interfaces. These require an ABAP team that knows APO internals.
- **Transparent Z-tables on the DB layer** — accessible via custom RFC/BAPI wrappers that Nestlé's ABAP team provides. SnapLogic's RFC Snap calls these directly.
- **BW/BW4HANA extractors** — if APO data is already extracted to BW, SnapLogic can read from BW via JDBC or OData.

SnapLogic's role: Once the extraction interface is confirmed and implemented by Nestlé's SAP/APO team or SI partner, SnapLogic handles everything downstream — transformation, routing, scheduling, and loading to Azure Data Lake. The Groundplex connects to whatever BAPI or DB interface is exposed.

Action item: Engage a certified SAP APO specialist or SI partner to confirm which tables are liveCache vs. DB-resident and what BAPI/extractor interfaces exist. This is a one-time discovery that unblocks the entire architecture.

CURRENT STATE (MARCH 12 DETAIL)

Sources: SharePoint (complex Nestlé sites, documents, file libraries), SmartFactory IoT devices (OPCUA/MQTT), PDFs, Word documents, videos, images, personal drives

Target: Azure DB Staging Layer (SharePoint migration), Azure Data Lake or Snowflake (IoT/unstructured)

Current problem: Each team connects to SharePoint/IoT sources independently, using their own tools (ADF, direct connectors) with no reusable strategy. No central metadata tagging, no visibility on what is being moved or duplicated.

Renan (March 12): *"We're doing it, but we're doing it without a strategy. People are just using connectors desperately."*

Scope for UC#3 (SharePoint migration): Data, attachments, comments, file library migration with column name changes, MVP with LLM-assisted pipeline creation

KEY SNAPLOGIC COMPONENTS

SharePoint Online Snap Pack

OPCUA Snap Pack

MQTT Snap Pack

SnapGPT Pipeline Gen

Mapper Snap

Azure SQL Insert

File Reader Snap

Binary Handler

Snowflake Bulk Load

SNAPLOGIC APPROACH

SharePoint: SharePoint Online Snap Pack provides native connectivity for content, metadata, file libraries, and list items. SnapGPT accelerates pipeline authoring for standard patterns — generating initial pipeline configurations from natural language, which engineers then review and extend. For complex migrations with column name changes and custom schema transformations, SnapGPT reduces the authoring time but does not replace pipeline design expertise. Column mapping with Mapper Snap handles name changes and schema drift.

IoT / SmartFactory (OPCUA/MQTT): SnapLogic natively connects to OPCUA endpoints and MQTT brokers via dedicated Snap Packs — collecting sensor data, machine logs, and IoT device streams from factory floors. Data is transformed and loaded to Snowflake or Azure Data Lake for analytics and AI consumption.

Governance layer: A shared SnapLogic project with approved, templated pipelines for each unstructured source type — teams use these rather than building independently. Role-based access ensures consumption is tracked and metadata is consistently applied.

VALUE DELIVERED

- SnapGPT replaces manual ADF/Logic App configuration for SharePoint
- Single pipeline handles data + attachments + comments + file metadata
- Reusable IoT ingestion template for OPCUA/MQTT factory devices
- Centralised unstructured data strategy — publish once, teams consume
- Audit trail for compliance during migration and ongoing ingestion
- Foundation for AI/LLM model consumption of unstructured data at scale

Unstructured Data Strategy — The Real Ask: Nestlé does not need SnapLogic to replace what teams are already doing ad-hoc. They need SnapLogic to provide the *published, reusable, governed pattern* that teams adopt instead of going rogue. Renan confirmed this explicitly: *"We should be implementing this strategy from us as an integration stream — publish a model and say, this is how you access SharePoint data. This is how you do it."* SnapLogic's Shared Component Library and project-level governance make this achievable out of the box.

4

Financial Data Replication Monitoring (Snowflake → ServiceNow)

USE CASE DETAILS

Source: Snowflake (financial replication data)

Target: ServiceNow (ticket enrichment + status updates)

Outputs: What happened, when, who is working on it, resolution time, prevention measures

SNAPLOGIC ARCHITECTURE USED

Snowflake Select Snap queries market-specific incidents → Mapper extracts incident/request numbers → ServiceNow REST API integration enriches with current status → Response assembled and delivered. The core pipeline pattern (Snowflake → ServiceNow enrichment loop) is baseline SnapLogic. Agent-based intelligent response generation (Snowflake Cortex, ServiceNow Assist) is an optional extension for teams with those capabilities enabled — not a prerequisite for the core use case.

This use case demonstrates how Sriven and Nestlé can rapidly deliver Snowflake-to-ServiceNow incident monitoring using SnapLogic's native connectors — without custom middleware. The pattern is reusable across Nestlé's operational monitoring landscape and becomes a template asset in the SnapLogic project space.

REQUIREMENTS — CURRENT ARCHITECTURE (MARCH 12 DETAIL)

Source: SAP S/4 cFIN (Finance)

Target: Snowflake

Current mechanism: **SAP SLT (Landscape Transformation Replication Server) + SNP Kyano Glue** add-on

SLT mechanism: DB triggers (INSERT/UPDATE/DELETE) on source tables → writes delta records to SLT logging table → SLT pushes to target. Trigger fires per row change, batches can be single records or thousands.

Current blocker: SNP Kyano Glue is required for CDS view extraction via SLT — it is a third-party add-on

Nestlé wants to eliminate

Volume: 100 MB/second data transfer

Frequency: Every minute replication

Operations: Insert only with complex filter criteria (RLDNR, ACCASTY, BTTYPER, etc.)

KEY SNAPLOGIC COMPONENTS

Ultra Pipeline (Headless)

SAP IDoc Listener

SAP RFC Snap

Filter Snap (complex criteria)

Snowflake Snowpipe Streaming

FeedMaster Node

SCD Type 2 Pattern

SNAPLOGIC APPROACH — KYANO GLUE REPLACEMENT

SnapLogic offers **two complementary approaches** to replace SNP Kyano Glue:

Option A — JDBC Direct to SLT Logging Table (Recommended)

SnapLogic's JDBC Snap Pack reads directly from SLT's internal logging table using a Groundplex on-prem Snapplex with database connectivity to the SAP HANA DB. Change records (INSERT/UPDATE/DELETE flags) are polled at configurable intervals, filtered per criteria (RLDNR, ACCASTY, BTTYPER), and streamed to Snowflake via Snowpipe. This approach keeps SLT's change capture mechanism intact while replacing Kyano Glue as the delivery layer.

Option B — IDoc/ABAP CDC via Ultra Pipeline

SAP IDoc Listener Snap in a Headless Ultra Pipeline continuously receives SAP-emitted change events without polling. Filter criteria are applied in the SnapLogic pipeline. Better for event-driven architectures where SLT is removed entirely.

PERFORMANCE ARCHITECTURE

- Multiple Ultra instances across Snapplex nodes (horizontal scale)
- FeedMaster for reliable message queuing + load balancing
- 100 MB/s sustained throughput (proven at Siemens)
- High-availability with no single point of failure
- Schema Registry validation for data quality

Architecture Note — Oracle DB Layer: Nestlé's SAP S/4 runs on **Oracle DB**, with database-level triggers already implementing CDC at the Oracle layer. This means SnapLogic's Oracle JDBC Snap (Groundplex on-prem) can read directly from Oracle's trigger-generated change log tables — no SLT setup, no IDoc configuration, no SAP-layer extraction dependency. This is the simplest and most reliable CDC path for these use cases. Confirm which Oracle change log tables are in use and verify JDBC read permissions for the Groundplex host.

Customer Reference — Siemens Master Data Distribution: SAP S/4HANA CDC at 100 MB/s using SnapLogic Ultra Pipelines + Kafka — every-minute replication with sub-second processing. High-availability architecture with FeedMaster nodes. **SNP Kyano Glue completely eliminated.** The Siemens architecture is the direct blueprint for Nestlé's streaming requirements.

REQUIREMENTS

Source: SAP S/4 HANA ERP

CDS View: MBV_MBEWH (plant stock valuation)

Target: Snowflake

Operations: Full delta replication — inserts, updates, deletes

Volume: 1 MB/second

Frequency: Every minute

KEY SNAPLOGIC COMPONENTS

SAP S/4HANA Read Custom Query

Ultra Pipeline (Headless)

Delta/Watermark Pattern

Router Snap (I/U/D)

Snowflake Merge/Upsert

Snowflake Execute

SNAPLOGIC APPROACH

SnapLogic's SAP S/4HANA Snap Pack provides **native CDS view connectivity** without SNP Kyano Glue. SAP S/4HANA Read Custom Query Snap accesses the MBV_MBEWH CDS view directly via OData/RFC. Delta tracking is handled by SnapLogic's incremental load patterns with watermark-based change detection. Insert/Update/Delete operations are routed to Snowflake using SCD Type 2 merge patterns via Snowflake Execute Snap.

BOTTLENECK RESOLUTION

- No SNP Kyano Glue license or setup required
- CDS view access via standard OData — no client-specific ABAP config
- Built-in I/U/D routing in pipeline logic
- 1 MB/s volume well within SnapLogic's capacity

Customer Reference — Ascential (FMCG): Delta replication (insert/update/delete) from SAP to Snowflake using Kafka for reliable streaming. High-volume data transactions (hundreds of millions GBP). Supports future IoT and Industry 4.0 initiatives. **Native SAP CDS view connectivity — no third-party add-ons required.**

REQUIREMENTS

Source: SAP S/4 HANA ERP

Middleware: Confluent (Kafka)

Target: Snowflake (merge into target table)

Operations: All changes — insert, update, delete

Volume: 100 MB/second

Frequency: Every minute (real-time)

SNAPLOGIC APPROACH

This is SnapLogic's **definitive sweet spot**. Native Kafka Consumer/Producer Snaps connect to Confluent with full Schema Registry (Avro + JSON Schema) support. Headless Ultra Pipelines continuously consume Kafka topics and merge into Snowflake via Snowpipe Streaming Snap. Change detection (I/U/D) captured at SAP via IDoc/ABAP CDC, published to Kafka topics, and merged into Snowflake in real-time.

KEY SNAPLOGIC COMPONENTS

Kafka Consumer Snap

Kafka Producer Snap

Confluent Schema Registry

Ultra Pipeline (Headless)

SAP IDoc Listener

Snowflake Snowpipe Streaming


Snowflake Merge

FeedMaster HA Node

PERFORMANCE SPECIFICATIONS

- 100 MB/s sustained streaming throughput (proven at Siemens)
- Sub-second processing latency via Ultra Pipelines
- Horizontal scaling across multiple Snaplex nodes
- Schema Registry validation ensures data quality at scale
- Retry mechanisms and error routing for reliability

SIEMENS-BLUEPRINT ARCHITECTURE: SAP → KAFKA → SNOWFLAKE REAL-TIME



Customer References:

Siemens: 100 MB/s SAP CDC via Kafka + Ultra Pipelines → Snowflake. Proven HA architecture with FeedMaster nodes. Exact Nestlé specification.

Wendy's (Food & Beverage): Mission-critical PoS data streaming at 10x Informatica performance. Real-time NCR→Internal Systems pipeline, demonstrating food industry streaming reliability.

Thayer Distribution (Food): 10x faster data movement for food distribution operations. Self-service for operational teams.

REQUIREMENTS

Source: SAP S/4 HANA ERP (ABAP Program ALV output)

Targets: Snowflake / SQL Server / SAP (multi-target)

Volume: 1 MB/second

Frequency: Every minute

Challenge: Final target routing and multi-target synchronization

SNAPLOGIC APPROACH

ABAP ALV output is extracted via SAP RFC Snap — but requires Nestlé's ABAP team to wrap the ALV program in a callable Function Module (ALV is a display framework, not a direct extraction interface; the RFC wrapper is a one-time ABAP development effort). Once the RFC interface exists, SnapLogic's Router Snap dynamically routes to Snowflake, SQL Server, or back to SAP based on data classification logic. Multi-target synchronization is handled through parallel Pipeline Execute Snaps. This directly supports Nestlé's requirement for flexible data routing without replicating pipeline logic per target environment.

KEY SNAPLOGIC COMPONENTS

SAP RFC Snap (ABAP FM)

Router Snap

Snowflake Insert

SQL Server Insert

SAP S/4HANA Write

Pipeline Execute (Parallel)

Ultra Pipeline

VALUE FOR NESTLÉ

- Single pipeline definition routes to all 3 targets
- Reusable ABAP extraction template across ALV programs
- Eliminates per-target custom code development
- Supports Nestlé's "no replicated logic" requirement

Sigma Framework Best Practice — Multi-Target Routing: Use SnapLogic's Router Snap with configurable destination accounts. Store target routing rules as pipeline parameters (not hardcoded), enabling the same pipeline to serve Snowflake, SQL, or SAP targets through configuration changes alone. This is the reusability pattern Nestlé cited as a core requirement.

SnapLogic Platform Architecture & Enterprise Capabilities

How SnapLogic's architecture addresses Nestlé's four core enterprise requirements: SAP connectivity, performance/scalability, governance & security, and lowest TCO.

Core Enterprise System Access — SAP

SnapLogic provides the deepest native SAP connectivity in the integration market — covering the full SAP ecosystem without third-party add-ons.

- ✓ SAP S/4HANA On-Premise Snap Pack (Read, Create, Update, Delete, Batch Read/Write, Custom Query)
- ✓ SAP S/4HANA Cloud Snap Pack (OData V2, full CRUD)
- ✓ SAP RFC/BAPI Execute, IDoc Read/Write/Listener
- ✓ SAP ECC 6.0 support via JCo Driver (3.1.x)
- ✓ CDS View extraction via OData/RFC without SNP Kyano Glue
- ✓ SSL + Basic Auth account types for secure SAP connectivity
- ✓ SAP APO, SAP Ariba, SAP HANA database support
- ✓ Groundplex on-premise for direct SAP LAN access, no VPN tunneling

Snowflake — Native Platform Integration

SnapLogic is a certified Snowflake partner with a purpose-built Snap Pack covering the full Snowflake data lifecycle.

- ✓ Snowflake Select, Insert, Update, Delete, Execute, Lookup
- ✓ Snowflake Bulk Load/Upsert for high-volume batch operations
- ✓ Snowflake Snowpipe Streaming — real-time continuous ingestion, configurable flush (1–60s)
- ✓ Vector Search support for AI/ML workloads
- ✓ Time Travel queries for historical data access
- ✓ Iceberg Tables support for partitioned analytics data
- ✓ Metadata caching for improved query performance
- ✓ RBAC and secure credential management via Vault integration

SAP APO Custom Table Access Methods — Decision Guide

SAP APO runs on the NetWeaver ABAP stack. Custom Z-tables can be accessed via three protocols in SnapLogic. The right choice depends on data type, volume, latency, and whether CDS views exist. This guide applies to UC #1 and UC #2.

RFC / BAPI

Primary Method for Z-Tables

Direct ABAP Function Module calls via SAP JCo connector. The **SAP RFC Snap** and **SAP BAPI Execute Snap** call standard or custom function modules to read Z-table data with full control over fields, filters, and pagination.

When to use:

- Batch extraction of APO Z-tables
- High-volume partitioned reads (UC #2)
- Complex ABAP transformation logic (UC #1)
- Tables with no published CDS/OData view

Key snaps:

SAP RFC Snap SAP BAPI Execute

APO liveCache note: Most APO planning tables reside in liveCache (in-memory), not in the relational DB layer. RFC_READ_TABLE does not apply. The extraction method must be confirmed by an SAP APO specialist — SnapLogic's RFC Snap calls whatever BAPI or DB interface Nestlé's team exposes.

IDoc (Intermediate Document)

Event-Driven / CDC

SAP emits IDocs on business events (goods movements, order creation, master data changes). SnapLogic's **SAP IDoc Listener Snap** continuously receives these via an always-on Ultra Pipeline — no polling required.

When to use:

- Real-time change capture (UC #5 CDC pattern)
- Order-driven or event-triggered integrations
- Master data distribution across systems
- *Not* suitable for bulk historical Z-table extraction

Key snaps:

SAP IDoc Listener SAP IDoc Read

SAP-side setup: IDoc partner profile + port configuration required in SAP WE20/WE21. Nestlé's Basis team handles this — typically 1-2 hours.

OData / CDS View

S/4HANA Standard Objects

CDS views published as OData services are accessed via SnapLogic's **SAP S/4HANA Read Custom Query Snap** — RESTful, stateless, no custom ABAP required. Standard S/4HANA objects (MBV_MBEWH, MATDOC, etc.) are available out of the box.

When to use:

- Standard S/4HANA objects with existing CDS views
- UC #4 and UC #6 (MBV_MBEWH, delta replication)
- Incremental load with \$filter / watermark pattern
- *Less applicable* to APO Z-tables — CDS views rarely published

Key snaps:

SAP S/4HANA Read Custom Query Snap

For APO Z-tables specifically:

Confirm with Nestlé's ABAP team whether CDS views exist. If not, default to RFC/BAPI.

Performance & Scalability — Ultra Pipelines

Ultra Pipelines are SnapLogic's high-performance execution engine designed for mission-critical, always-on integrations at enterprise scale.

- ✓ Always-on execution — eliminates startup overhead for real-time requirements
- ✓ 100 MB/s+ sustained CDC throughput (proven at Siemens)
- ✓ Sub-second processing latency for streaming use cases
- ✓ Headless Ultra Tasks for continuous Kafka consumption
- ✓ Low-latency Ultra Tasks for real-time API query responses
- ✓ FeedMaster Nodes for reliable message queuing and load balancing
- ✓ Horizontal scaling across multiple Snaplex execution nodes
- ✓ High-availability architecture with no single point of failure
- ✓ Configurable batch sizes per Snap for optimal throughput tuning

Hybrid Architecture — Groundplex & Cloudplex

SnapLogic's hybrid execution model meets Nestlé's security and network requirements while providing cloud elasticity where needed.

- ✓ Groundplex: On-premise Snaplex for direct SAP LAN access, zero data leaves the network
- ✓ Cloudplex: SnapLogic-managed cloud for external API and SaaS integrations
- ✓ Hybrid pipelines: split execution between on-prem (SAP) and cloud (Snowflake) as needed
- ✓ No general internet access required on Groundplex — targeted outbound rules only
- ✓ Groundplex connects to SnapLogic Control Plane via port 443 only
- ✓ External service access (Salesforce, APIs) via specific firewall rules — not general internet
- ✓ On-prem execution meets Nestlé compliance and data residency requirements
- ✓ Fully managed Cloudplex for cloud-to-cloud integrations (zero infrastructure management)

Governance & Security

Enterprise-grade governance built into the platform — not bolted on. Supports Nestlé's requirement for Citizen Developer enablement with proper oversight.

- ✓ Role-Based Access Control (RBAC) at organization, project, and asset level
- ✓ 90-day audit logging — all user actions, pipeline executions, configuration changes
- ✓ Centralized credential management — HashiCorp Vault, AWS Secrets Manager integration
- ✓ Pipeline validation rules enforced at project and org level
- ✓ Approval workflows for pipeline promotion to production
- ✓ API policy enforcement (auth, rate limiting, IP allowlisting)
- ✓ SSL/TLS encryption for all SAP, Snowflake, and cloud connections
- ✓ Shared Component Library for standardized, approved pipeline patterns
- ✓ OpenTelemetry-compatible monitoring for operational observability

Citizen Developer & Lowest TCO

SnapLogic is designed to reduce dependency on specialized integration engineers — directly addressing Nestlé's 4-6 day development time and consolidation goals.

- ✓ SnapGPT: AI copilot generates pipeline configurations from natural language
- ✓ AutoSync: No-code endpoint-to-endpoint data synchronization for standard patterns
- ✓ 1000+ pre-built Snaps — no custom code for common integration patterns
- ✓ Visual drag-and-drop Designer — accessible to business users and analysts
- ✓ Template library for reusable, approved integration patterns
- ✓ Role-based self-service with governance guardrails
- ✓ Eliminates ADF + BODS + SNP Kyano Glue + custom scripts licensing
- ✓ Unified platform: batch, real-time, API, streaming, unstructured — all in one
- ✓ Reduces integration development time from 4-6 days to hours (per Mikron reference)

SnapLogic vs. Current Tools — Why Consolidate

Nestlé's current stack includes ADF, MuleSoft, BODS, SNP Kyano Glue, and custom scripts. This is the honest comparison.

Capability	SnapLogic	Azure ADF	MuleSoft	SAP BODS
Primary strength	Unified batch + streaming + API + unstructured data — one platform	Batch data movement, Azure-native workloads	API management, transactional integration, B2B	SAP-to-SAP ETL, BusinessObjects data services
SAP connectivity	Native RFC, IDoc, OData, JDBC — covers full SAP stack. SAP APO extraction requires SI partner for liveCache objects	SAP connector available; complex SAP patterns require custom linked services or SAP BW connectors	SAP connector available; deep SAP requires MuleSoft SI partner expertise	Deep SAP-native; best suited for SAP-to-SAP; limited non-SAP target support
Real-time / CDC	Ultra Pipelines, Kafka Consumer/Producer, Snowpipe Streaming — proven 100 MB/s. Oracle DB trigger CDC via JDBC	Limited native CDC; requires Azure Event Hub or Logic Apps for real-time patterns	Streaming via Anypoint MQ; Kafka connector available but separate licensing tier	Batch-oriented; real-time requires additional SAP components (SLT, Event Mesh)
AI-assisted development	SnapGPT: natural language pipeline generation, AutoSync no-code sync	Copilot in Azure Data Factory (limited; primarily ADF-native patterns)	Copilot capabilities in Anypoint; focused on API design	No AI-assisted pipeline authoring
Unstructured data	SharePoint, OPCUA, MQTT, PDFs, IoT streams — native Snap Packs	SharePoint connector; limited IoT/unstructured handling without Logic Apps	Not a core use case; typically requires custom connectors	Not designed for unstructured data sources
Snowflake integration	Certified Snowflake partner; bulk load, Snowpipe Streaming, Iceberg, Vector Search	Snowflake connector available; no native Snowpipe Streaming	Snowflake connector via Anypoint; no native Snowpipe	Snowflake target supported via JDBC; no streaming capability

Capability	SnapLogic	Azure ADF	MuleSoft	SAP BODS
On-premise execution	Groundplex: on-prem Snaplex for direct SAP LAN access, no VPN tunneling	Self-hosted Integration Runtime (IR) — comparable on-prem execution	On-prem Mule Runtime available but requires infrastructure management	On-premise only; no cloud-native option without SAP Data Intelligence
Consolidation opportunity for Nestlé	Replaces ADF (batch + CDC), BODS (ETL), SNP Kyano Glue (CDC add-on), and reduces custom script sprawl in one platform	Replaces batch pipelines; does not address MuleSoft/BODS/Kyano scope	Replaces API layer; does not address bulk data, CDC, or unstructured scope	Replaces only SAP ETL; no path to real-time, unstructured, or non-SAP targets
SAP implementation expertise	Platform capability is strong. Deep SAP implementation (APO liveCache, complex ABAP) requires certified SI partner	Microsoft SI ecosystem; SAP-specific expertise via SAP-certified ADF partners	Large SI partner ecosystem; SAP + MuleSoft is common but expensive	SAP-native — but requires SAP expertise in-house or via SAP-certified partner

Delivery Model — Sriven × Nestlé Tag-Team

SnapLogic provides the platform, the connectivity, the pipeline runtime, and the AI-assisted development tooling. **Sriven**, as a certified SnapLogic implementation partner, brings the hands-on delivery expertise to design, build, and validate pipelines against Nestlé's environment. **Nestlé's SAP team** brings the system architecture knowledge — ABAP extraction interfaces, APO liveCache structures, S/4HANA configuration — that no external party can replicate.

This is the intended operating model: SnapLogic as platform, Sriven as implementation lead, Nestlé as SAP authority. The sandbox validation phase will allow Sriven and Nestlé to jointly confirm each use case approach before production commitment.

Customer Reference Portfolio

Verified customer implementations matching Nestlé's use cases, industries, and technical requirements.

Customer	Industry	Relevant Use Case(s)	Match Criteria	Key Outcome
Siemens Master Data Distribution	Manufacturing	UC #5 (CDC), UC #7 (Kafka Streaming)	100 MB/s CDC, Ultra Pipelines, Kafka, every- minute replication, SAP S/4HANA source, Snowflake target	Proven 100 MB/s sustained throughput. Sub-second latency. HA architecture. Direct blueprint for Nestlé's streaming requirements.
Mikron	Food Manufacturing	UC #2 (SAP APO → Snowflake)	High-volume SAP-to- Snowflake, partitioned queries, on-prem Snaplex, 3 medium nodes	Integration time: 2-3 days → 2 hours (75% reduction). Team: 3-5 → 1-2 people. Unified SAP, Salesforce, Snowflake platform.
Nielsen	Consumer Packaged Goods	UC #2 (SAP Migration)	SAP data migration at CPG scale, tight timelines, elimination of external consultants	Met aggressive Q1 SAP go-live. Internal teams manage all SAP integrations. Zero external consultants post- implementation.
Ascential	FMCG	UC #6 (CDS View Delta Replication)	Delta replication (I/U/D), Kafka messaging, SAP- to-Snowflake, FMCG scale	Complex SAP integrations at FMCG scale. High-volume CDC with hundreds of millions GBP in transactions. IoT-ready architecture.
Wendy's	Food & Beverage	UC #7 (Real- Time Streaming)	Mission-critical real-time streaming, food industry, Informatica replacement	10x faster data movement vs. Informatica. Improved reliability for critical PoS data. Self- service for operational teams.
Thayer Distribution	Food Distribution	UC #7 (High- Volume Streaming)	High-volume food industry data movement, self-service platform, operational efficiency	10x faster data movement. Improved profitability. Self- service integration for business users.
General Mills	FMCG / Food	All SAP + Data Platform UCs	FMCG enterprise, SAP core systems, data platform consolidation	Enterprise integration consolidation at global FMCG scale. SAP + cloud data platform connectivity with governance.

S

Siemens

Manufacturing — Use Cases
5 & 7

SAP Master Data
Distribution at 100
MB/s via Ultra
Pipelines + Kafka. The
definitive reference
architecture for
Nestlé's streaming and
CDC requirements.

100 MB/s Proven

M

Mikron

Food Manufacturing — Use
Case 2

Replaced ADF for SAP
APO → Snowflake.
Native partition
handling for multi-
million row tables on
Groundplex with 75%
faster delivery.

75% Faster Integration

W

Wendy's

Food & Beverage — Use
Case 7

Mission-critical real-
time PoS data
streaming. Replaced
failing Informatica with
SnapLogic, achieving
10x performance for
food industry
operations.

**10x Performance vs
Informatica**

Summary

Nestlé's integration environment is complex by design: a global SAP estate spanning APO, S/4HANA, and cFIN; a strategic commitment to Snowflake as the analytics platform; a declared 2026 priority around unstructured data; and a recognised need to move away from a fragmented toolset (ADF, BODS, SNP Kyano Glue, custom scripts) that creates bottlenecks and dependency on specialist developers.

This document has assessed eight identified use cases against SnapLogic's platform capabilities and verified customer implementations. Based on the shared understanding developed between SnapLogic, Sriven, and Nestlé, the platform is expected to address the full breadth of integration requirements from a single platform — high-volume batch, real-time CDC at 100 MB/s, Kafka-based streaming, CDS view replication, SharePoint and IoT ingestion, and AI-assisted pipeline development — without third-party add-ons like SNP Kyano Glue.

Validation against Nestlé's sandbox will confirm the approach for each use case before any production commitment.

Use Case	Current Tool(s) Replaced	SnapLogic Fit	Key Proof Point
UC #1 — SAP APO → OMP / Azure Data Lake	SAP BODS (downstream ETL)	• VALIDATED	RFC/BAPI + Groundplex on-prem. ABAP extraction design confirmed by Nestlé's SAP team or SI partner.
UC #2 — SAP APO → Snowflake STG (Priority #1)	Azure Data Factory (ADF)	• DIRECT BLUEPRINT	Mikron: 2-3 days → 2 hours (75% reduction). Partitioned SAP → Snowflake at scale.
UC #3 — Unstructured Data Strategy (SharePoint, IoT, docs)	Ad-hoc ADF / direct connectors (no central strategy)	• VALIDATED	SharePoint Online, OPCUA/MQTT Snap Packs. SnapGPT-assisted pipeline generation for governed, reusable templates.
UC #4 — Financial Replication Monitoring (Snowflake → ServiceNow)	Custom scripts / manual process	• VALIDATED	Native Snowflake + ServiceNow connectivity. Pattern reusable across operational monitoring use cases.
UC #5 — High-Volume CDC (S/4 cFIN → Snowflake, 100 MB/s)	SNP Kyano Glue + SLT (add-on eliminated)	• DIRECT BLUEPRINT	Siemens: 100 MB/s SAP CDC via Ultra Pipelines + Kafka. Oracle DB JDBC path simplifies further.
UC #6 — CDS View Delta Replication (MBV_MBEWH → Snowflake)	SNP Kyano Glue (eliminated)	• NATIVE CAPABILITY	Standard OData access to CDS views — no third-party add-on. Watermark-based delta pattern.

Use Case	Current Tool(s) Replaced	SnapLogic Fit	Key Proof Point
UC #7 — SAP S/4 → Kafka (Confluent) → Snowflake, 100 MB/s	ADF / custom pipeline	DIRECT BLUEPRINT	Siemens architecture. Kafka Consumer/Producer, Ultra Pipelines, Snowpipe Streaming — production-proven.
UC #8 — ABAP Extractions → Multi-Target (Snowflake / SQL / SAP)	Per-target custom scripts	VALIDATED	Single parameterised pipeline with Router Snap handles all 3 targets. Reusable across ALV programs.

Platform Consolidation — One Tool, Full Stack

SnapLogic replaces ADF (batch + CDC), SAP BODS (ETL post-extraction), and SNP Kyano Glue (CDC add-on) in a single licensed platform. The unstructured data and streaming scope — currently covered by ad-hoc tools or not addressed at all — is included without additional point solutions.

Development Velocity — From Days to Hours

Nestlé's confirmed 4-6 day integration development cycle is the target. SnapGPT-driven pipeline generation, 1000+ pre-built Snaps, and a reusable template library address this directly. Mikron reduced equivalent SAP-Snowflake development from 2-3 days to 2 hours. Nielsen eliminated external consultants post-go-live.

Architecture Readiness — SAP Complexity Is Acknowledged

SAP APO liveCache extraction, Oracle CDC, and ABAP ALV wrapping all require Nestlé's SAP/ABAP team or an SI partner's expertise to define the extraction interface. SnapLogic owns everything downstream: transformation, routing, scheduling, Snowflake loading, and monitoring. The boundary is clear, and it applies to every integration platform on the market.

The case for SnapLogic at Nestlé: The Siemens architecture is the streaming blueprint. Mikron is the SAP-Snowflake blueprint. Nielsen demonstrates the path to internal team ownership. All three apply directly to Nestlé's stated priorities. SnapLogic consolidates the toolchain, reduces development time, and provides a governed foundation for the unstructured data strategy — without requiring Nestlé to replace MuleSoft, which continues to serve the API and transactional integration layer where it is already embedded.

The path forward: Sriven and Nestlé validate the proposed approach in the provisioned sandbox, starting with UC #2. Each use case tested in sandbox becomes a confirmed delivery commitment. SnapLogic supports the process throughout.

SnapLogic Platform Architecture & Glossary

How SnapLogic is structured — so Nestlé and Sriven speak the same language from day one of the sandbox

High-Level Architecture

CONTROL PLANE (SNAPLOGIC CLOUD)

- **Designer** — browser-based pipeline IDE
- **Manager** — monitoring, scheduling, access control
- **SnapLogic IIP** — Intelligent Integration Platform core
- **SnapGPT** — AI pipeline generation layer
- **Metadata & Repository** — pipeline storage, versioning

DATA PLANE (CUSTOMER-CONTROLLED)

- **Groundplex** — on-premises or private cloud runtime
- **Cloudplex** — SnapLogic-managed cloud runtime
- **Ultra Pipelines** — streaming / high-throughput path
- **Snaplex nodes** — JVM worker processes (auto-scaled)
- **Local File System / JDBC / SAP JCo** — connector layer

Key principle: All data processing happens inside the Data Plane — data never transits through the SnapLogic Control Plane. The Control Plane only carries pipeline metadata, configuration, and monitoring signals. This means Nestlé's SAP and Snowflake data stays within Nestlé's network perimeter when using a Groundplex.

Recommended Deployment Topology for Nestlé

On-Prem Groundplex

Deployed inside Nestlé network. Direct JDBC to Oracle/SAP DB, SAP JCo for RFC/IDoc, file system access. Data never leaves the perimeter.

Cloudplex (optional)

For cloud-only integrations (Snowflake ↔ SaaS). SnapLogic-managed, zero infrastructure burden on Nestlé. Can run in parallel with Groundplex.

Ultra Pipelines

High-throughput streaming mode (100 MB/s+). Used for SAP CDC → Kafka → Snowflake (UC #7). Low-latency, continuous data flows.

SnapLogic Terminology Glossary

Reference for Nestlé stakeholders encountering SnapLogic terminology for the first time. Terms are grouped by layer.

CORE PLATFORM COMPONENTS

Term	Definition
Snap	The atomic unit of integration logic. Each Snap performs one operation — read from a database, transform a field, write to Snowflake, call an API. Snaps are pre-built, configured via a property panel, and connected visually. SnapLogic ships 1000+ Snaps out of the box.
Snap Pack	A versioned bundle of Snaps for a specific technology (e.g. SAP Snap Pack, Snowflake Snap Pack, Kafka Snap Pack). Snap Packs are independently versioned and upgradeable without affecting other integrations.
Pipeline	A visual workflow of connected Snaps that defines an integration flow — source → transform → target. Pipelines are the deployable artefact in SnapLogic. They can be triggered on a schedule, by an API call, by a file event, or invoked by another pipeline.
Ultra Pipeline	A specialised pipeline mode for high-throughput, low-latency streaming. Processes data continuously rather than in batches. Used for Kafka consumers/producers and real-time CDC scenarios. The Siemens 100 MB/s SAP→Snowflake flow runs via Ultra Pipelines.
Task	A scheduled or triggered execution of a pipeline. Tasks define when a pipeline runs, on which Snaplex, with what runtime parameters. Tasks are managed from the Manager UI or via REST API.
Project	An organisational folder in SnapLogic that groups pipelines, accounts, files, and tasks. Projects map to teams, environments, or business domains. Access control and permissions are applied at project level.
Project Space	A top-level namespace that contains multiple Projects. Typically represents an organisation or major division. In multi-tenant deployments, each Project Space is isolated.
Account	A reusable credential or connection configuration (e.g. "SAP Production Account", "Snowflake Data Lake Account"). Accounts are defined once and referenced by many pipelines. Credentials are stored encrypted and never appear in pipeline code.

Term	Definition
Snaplex	The execution engine that runs pipelines. A Snaplex is a cluster of one or more JVM nodes. It connects outbound to the SnapLogic Control Plane for metadata and pipeline definitions, but all data processing happens locally. There are two types: Groundplex and Cloudplex.
Groundplex	A Snaplex deployed inside the customer's infrastructure (on-premises, private cloud, or VPC). The customer installs and manages the JVM nodes. Required for any integration where data must not leave the network — SAP on-prem, Oracle on-prem, sensitive Nestlé datasets.
Cloudplex	A Snaplex hosted and managed by SnapLogic in the public cloud. No infrastructure to provision or maintain. Suitable for cloud-to-cloud integrations (e.g. Salesforce → Snowflake). Not appropriate where data must stay on-prem.
Node	An individual JVM process within a Snaplex. Multiple nodes provide horizontal scale and high availability. Pipelines are distributed across nodes automatically. Adding nodes increases throughput without reconfiguring pipelines.
Control Plane	The SnapLogic-hosted SaaS layer. Hosts the Designer IDE, Manager console, pipeline metadata repository, SnapGPT, and monitoring dashboards. Does not process customer data — only pipeline definitions, configurations, and telemetry.
Data Plane	Where actual data processing happens — the Snaplex nodes. Entirely within the customer's environment for Groundplex deployments. Separated from the Control Plane by design to ensure data residency and compliance.

Term	Definition
Designer	The browser-based IDE for building pipelines. Drag-and-drop Snaps onto a canvas, connect them, and configure properties. No code required for standard integrations. Supports inline expression language (JSONata-like) for field mapping and transformations.
SnapGPT	SnapLogic's embedded generative AI assistant. Accepts natural-language descriptions of an integration and generates a pipeline draft — correct Snaps selected, connected, and partially configured. Reduces initial pipeline build time from hours to minutes. Operates inside Designer.
Expression Language	The inline scripting syntax used inside Snap property fields for dynamic values, field mapping, and conditional logic. Based on a JSONata-derived syntax. Used for field-level transformation without needing a separate Script Snap.
Script Snap	A Snap that executes custom JavaScript or Python code inline within a pipeline. Used for complex transformation logic that exceeds what the Expression Language can handle. Runs inside the Snaplex JVM — no external runtime required.
Mapper Snap	A dedicated field-mapping Snap with a visual mapping grid. Maps source document fields to target document fields with optional transformations. Auto-suggests mappings when source and target schemas are available.
Router Snap	Splits a data stream into multiple branches based on conditions. Each branch can route to a different target system. Used in UC #8 to send ABAP extraction output to Snowflake, SQL Server, or SAP depending on payload type — eliminating per-target custom scripts.
Pipeline Parameters	Named variables passed into a pipeline at runtime. Enable a single parameterised pipeline to serve multiple source tables, environments, or configurations. Critical for the "build once, run many" reuse pattern central to the proposed Nestlé architecture.

INTEGRATION PATTERNS & SAP-SPECIFIC TERMS

Term	Definition
SAP Snap Pack	SnapLogic's native SAP connector bundle. Includes Snaps for RFC/BAPI calls, IDoc send/receive, OData (CDS Views), JDBC (SAP HANA / SAP DB), SAP JCo (Java Connector) for direct ABAP function module calls. Used across UC #1–#8 in this assessment.
SAP JCo	SAP Java Connector — the SAP-provided library that enables direct RFC/BAPI communication from Java applications to SAP ABAP systems. SnapLogic's SAP Snap Pack wraps JCo. Deployed on the Groundplex node inside Nestlé's network for secure on-prem SAP access.
CDC (Change Data Capture)	A data replication pattern that captures only rows that have changed since the last extraction — inserts, updates, deletes — rather than re-reading the full table. SnapLogic supports CDC via database log-based connectors (Oracle LogMiner, SAP SLT trigger-based, watermark columns on CDS Views).
Watermark	A stored high-water-mark value (typically a timestamp or sequence number) that marks the point up to which data has been successfully extracted. On the next run, the pipeline reads only records newer than the watermark. Used in the CDS View delta replication pattern (UC #6).
ABAP ALV	SAP ABAP List Viewer — a standard SAP reporting framework. In the context of UC #8, Nestlé's ABAP team wraps complex APO liveCache extractions in ALV-format ABAP programs that expose the data via RFC. SnapLogic calls the RFC and handles everything downstream.
IDoc (Intermediate Document)	SAP's native asynchronous message format for inter-system data exchange (e.g. SAP ERP → SAP WM, SAP → external WMS). SnapLogic's IDoc Snap Pack can send and receive IDocs, enabling event-driven warehouse management flows (UC #3).
OData / CDS View	Core Data Services (CDS) views are SAP's semantic layer on top of database tables, exposable as OData REST endpoints. SnapLogic reads CDS views via standard HTTP OData calls — no additional SAP add-on (such as SNP Kyano Glue) required. Used in UC #6 for delta replication.
Snowpipe Streaming	Snowflake's low-latency continuous ingestion API. SnapLogic's Snowflake Snap Pack supports Snowpipe Streaming for real-time data load — used in the UC #7 SAP → Kafka → Snowflake architecture to achieve sub-second latency into Snowflake tables.
Citizen Integrator	A business user or non-specialist developer who builds integrations using SnapLogic's low-code/no-code tools (Designer + SnapGPT) without deep programming knowledge. SnapLogic's federated model targets CoE-governed platforms where domain teams own their own integrations — the 3-year maturity model outlined for Nestlé.