STRATEGIC ALIGNMENT DOCUMENT — CONFIDENTIAL

Intelligent Integration for Choco's Scale

From n8n Migration to Enterprise iPaaS — ERP Orchestration, AI Readiness & the Platform for What Comes Next

- n8n Migration & Path Forward
- 12M+ Executions / Month
- AI-Ready Agentic Integration
- 150+ ERP Connectivity
- EU Data Residency & GDPR
- Food Commerce Scale

PREPARED FOR
Choco — Engineering & Operations Leadership

PREPARED BY
Konstantin Riegel, SnapLogic EMEA

DATE
March 2026

CLASSIFICATION
Confidential

Table of Contents

Strategic Alignment — SnapLogic × Choco | March 2026

01 — Management Summary

Choco has built one of the world's most ambitious platforms for B2B food commerce — connecting over 100,000 restaurants with thousands of food distributors across Europe and North America. To sustain that growth, Choco's integration infrastructure needs to do what n8n was never designed to do: operate reliably at 12M+ executions per month, connect 150+ distinct ERP systems, and serve as the backbone of a platform that food businesses depend on every day.

12M+	1,500+	150+	100k+
EXECUTIONS / MONTH	ACTIVE WORKFLOWS	ERP SYSTEMS CONNECTED	RESTAURANTS ON PLATFORM

THE MIGRATION MOMENT

n8n was the right choice at Choco's earlier stage — fast to deploy, developer-friendly, open source. At 12M executions/month and 1,500 workflows, the engineering cost of running n8n at scale now exceeds the cost of an enterprise platform. The migration moment has arrived.

ERP CONNECTIVITY AT SCALE

Choco's platform depends on integrating with 150+ distinct ERP systems used by food distributors globally. SnapLogic's 800+ pre-built connectors — including native ERP Snaps for SAP, Oracle, Microsoft Dynamics, Sage, and dozens more — eliminate the custom code that n8n requires for each new integration.

EU-FIRST BY DESIGN

Headquartered in Berlin, Choco operates under EU data regulations. SnapLogic's Groundplex hybrid deployment model allows processing to run within EU boundaries, with EU-region cloud hosting and GDPR-compliant data handling — not an afterthought, but a deployment option built into the platform.

AI READINESS

Choco's engineering team is building AI-augmented workflows on top of its integration layer. SnapLogic's SnapGPT accelerates pipeline development by 65%+, while the MCP Server integration model allows AI agents to invoke SnapLogic pipelines as tools — enabling the agentic automation layer Choco will need in 2026–2027.

OUR RECOMMENDATION

Before any full migration engagement, a **Proof of Value (POV)** is the right first step — limited in scope to **1–2 use cases**, supported by SnapLogic Solutions Engineering, and non-production. The POV's purpose is to prove that SnapLogic can drive additional value for Choco. If successful, it leads directly to a buying decision and full migration programme. A proper **scoping and sizing workshop** is required first so Choco receives the correct commercial offer for the engagement package that matches its objectives.

02 — Choco — Business Context & Strategic Objectives

Understanding Choco's business model, growth trajectory, and technical architecture is the foundation for any meaningful integration strategy.

🏢 Company Overview

Choco was founded in Berlin in 2018 to digitize the global food supply chain — starting with how restaurants communicate with food distributors. Before Choco, ordering ran on phone calls, faxes, and paper despite the sector's hundreds of billions in annual transaction volume.

- Founded: 2018, Berlin, Germany
- Founders: Daniel Khachab (CEO), Rogier Richter, Christopher Zemina, Julian Hammer
- Headquarters: Berlin, with offices in New York, Chicago, Paris, Madrid, Amsterdam
- Funding: Over $300M raised (Bessemer Venture Partners, Coatue, and others)
- Business model: SaaS platform connecting restaurants and food distributors via mobile-first ordering

⚙️ Integration Profile (Current)

Every restaurant order must flow to the corresponding distributor's ERP in real time, reliably, with full error traceability. The integration layer is the core of the product.

- **12M+ executions/month** — current throughput on n8n
- **1,500+ workflows** — active automations in production
- **150+ distinct ERP systems** — each distributor may run a different ERP
- Primarily REST/webhook-based inbound flows
- High variance in ERP API quality across the distributor base
- EU data residency requirements across all European operations

📈 Scale & Reach

Choco has grown from a Berlin-based startup to one of the most significant B2B food commerce platforms in the world, with a presence across multiple continents.

- 100,000+ restaurants and food businesses on the platform
- 20,000+ food distributors and suppliers connected
- Operations across Europe and North America
- Billions in gross merchandise value processed annually
- Average distributor reduces order processing time by 50%+

🎯 Strategic Objectives for 2026

Based on the discussions with the Choco engineering and operations leadership, the following strategic objectives are driving the platform modernisation initiative:

- Migrate from n8n to an enterprise-grade integration platform capable of sustaining growth to 50M+ executions/month
- Reduce the engineering overhead of onboarding new distributor ERP connections
- Achieve enterprise-grade reliability, observability, and error handling
- Maintain EU data residency compliance as the platform expands
- Build toward AI-augmented integration workflows for predictive operations
- Reduce time-to-onboard for new food distributors from weeks to days

WHY INTEGRATION IS CHOCO'S CORE PRODUCT RISK

For Choco, the integration layer is the product. Every order placed through the Choco app must reach the distributor's ERP in the right format, at the right time, with full traceability. A failed integration means a missed delivery. A delayed integration means a chef without ingredients at service prep. This is not an IT tool — it directly determines customer experience and commercial success.

03 — Food Commerce Industry: The Digitization Imperative

The food service supply chain is undergoing structural transformation, driven by the shift to digital ordering and demand for scalable integration across a fragmented ERP landscape.

$700B+	<30%	150+
GLOBAL FOOD SERVICE SUPPLY CHAIN (ANNUAL GMV)	DIGITIZED ORDERS BEFORE PLATFORMS LIKE CHOCO ENTERED MARKET	DIFFERENT ERP SYSTEMS IN ACTIVE USE ACROSS FOOD DISTRIBUTORS

📱 The B2B Ordering Transformation

Food service has been one of the last major B2B sectors to go digital. Phone calls, fax orders, and handwritten delivery notes persist because of extreme fragmentation: hundreds of thousands of restaurants and tens of thousands of distributors, each with their own ERP and processes.

- Platforms like Choco replace phone/fax with structured digital orders
- Structured orders must connect to distributor ERPs in real time
- Each distributor integration is custom — no industry-wide API standard exists
- As Choco scales, the integration challenge scales non-linearly

🔗 The ERP Connectivity Challenge

Food distributors run a highly fragmented ERP landscape — from bespoke food-industry systems (Produce Pro, BlueCart) to SAP S/4HANA and Oracle. Each integration is unique, with no industry-wide API standard.

- No standard API protocol across the distributor ERP market
- Mix of REST, SOAP, FTP/SFTP, EDI, and proprietary protocols
- Data model differences: order structures vary per ERP
- Reliability requirements: missed order = missed delivery = lost customer
- 150+ active ERP connections at Choco today; growing as platform expands

🌍 Industry Trends Shaping Integration Requirements (2025–2027)

Real-Time Order Visibility

Restaurants expect real-time confirmation of order receipt, not batch updates. This shifts integration from ETL patterns to event-driven, streaming architectures.

Predictive Inventory & Demand

Leading distributors are beginning to offer predictive ordering, which requires AI models consuming real-time order data — demanding an AI-ready integration layer that feeds ML pipelines.

Regulatory Pressure

EU food safety regulations, supply chain traceability requirements (EUDR), and GDPR combine to mandate full audit trails from order placement to delivery — which requires an integration platform with built-in observability.

04 — n8n at the Crossroads — Where It Excels, Where It Breaks

n8n is an excellent tool for its intended purpose. This section is an honest assessment of where n8n's architecture creates compounding costs at Choco's current scale — not because n8n is bad, but because Choco has grown beyond what it was designed to sustain.

✅ Where n8n Has Served Choco Well

- Fast to deploy — low setup friction for early-stage teams
- Developer-friendly — strong JavaScript/TypeScript support for custom logic
- Open source — no vendor lock-in at initial adoption stage
- Good UI for building and debugging individual workflows
- Flexible trigger system — webhooks, schedules, events
- Cost-effective at low-to-medium execution volumes

⚠️ Where n8n Creates Compounding Costs at Scale

- No native multi-tenant architecture — isolating 150+ ERP tenants requires custom infrastructure work
- Horizontal scaling requires significant custom setup (Kubernetes, queue workers, DB sharding)
- No built-in RBAC or SSO/SAML — enterprise access control requires workarounds
- Limited observability — no built-in business dashboards; monitoring requires external tooling
- No native circuit breakers or enterprise-grade retry policies
- SOC 2 / ISO 27001 compliance posture only on paid Cloud tier; self-hosted requires own programme

⚠️ **The 12M Execution Threshold:** n8n was designed for task automation at developer scale, not multi-tenant SaaS platforms processing millions of events per month. At Choco's volume, the engineering cost of maintaining n8n stability — custom scaling, bespoke monitoring, tenant isolation — has likely crossed the cost of an enterprise iPaaS.

The n8n Scaling Wall — What It Looks Like in Practice

Challenge	n8n Behaviour	Engineering Cost
Horizontal scale-out	Requires manual worker queue configuration, external Redis/PostgreSQL, custom Kubernetes setup	DevOps time: 2–4 weeks per scaling event; ongoing maintenance burden
Tenant isolation (150+ ERPs)	No native concept of tenants — isolation achieved via separate namespaces, manual tagging, or separate instances	Multiplied infrastructure cost; risk of cross-tenant data leakage without careful engineering
Error handling at 12M/month	Error handling is per-workflow, manual, inconsistent unless developers enforce standards	Growing technical debt; silent failures at scale without external monitoring
ERP connector maintenance	Custom code nodes required for complex ERP protocols (EDI, SOAP, SuiteQL, etc.)	Each new ERP integration is a development project; no reusable library of certified connectors

05 — SnapLogic Platform — Why It Fits Choco

SnapLogic is an enterprise-grade, AI-powered iPaaS built for organisations processing high volumes of data across complex, heterogeneous system landscapes. This section maps its capabilities directly to Choco's requirements.

🌊 Streaming Architecture — No Memory Ceilings

SnapLogic processes data as a continuous stream of documents, not as batch loads. Data flows snap-to-snap in parallel channels — meaning Choco's large ERP payload files (potentially millions of order records) are never loaded entirely into memory. This is the architectural reason SnapLogic scales where n8n hits limits.

- 100K–500K records/minute to Snowflake via native COPY INTO
- Parallel processing across pipeline branches automatically
- Execution scales elastically — burst workloads handled without pre-configuration

🔌 800+ Pre-built Snaps — Including ERP-Native

SnapLogic's Snap catalog covers virtually every system Choco's distributors run. Unlike n8n community nodes, SnapLogic Snaps are certified, version-managed, and maintained by SnapLogic engineering — not by individual developers who may abandon a project.

- SAP (S/4HANA, ECC, Business One) — native BAPI/RFC/OData support
- Oracle ERP Cloud, NetSuite (SuiteQL native)
- Microsoft Dynamics 365, Business Central, AX
- Sage, Infor, Epicor, Acumatica, and 30+ more ERP systems
- REST, SOAP, GraphQL, EDI, FTP/SFTP — all protocol-native

🏢 Multi-Tenant by Design

SnapLogic's SnapLogic Elastic Integration Cloud is built as a multi-tenant platform at the infrastructure level — not bolted on via configuration. For Choco, this means each food distributor ERP connection can be isolated, governed, and monitored as a first-class tenant without custom infrastructure work.

- Org-level isolation — each integration environment is fully separated
- RBAC with fine-grained permissions — per project, per pipeline, per environment
- SSO/SAML/OIDC built in — no workarounds required
- Audit logs at every layer — user actions, pipeline executions, data access

📊 Enterprise Observability — Built In

SnapLogic's monitoring capabilities are production-grade out of the box — no external APM tools required to understand what is happening across 1,500+ workflows and 12M+ executions/month.

- Real-time execution dashboard with per-pipeline throughput metrics
- Document-level error tracking — every failed record captured with reason
- SLA alerting — notify on latency, failure rate, or throughput thresholds
- Webhook, email, Slack, and Teams notification routing for error events

MIGRATION PATH — NOT A RIP-AND-REPLACE

SnapLogic's migration is a parallel-run model — n8n workflows continue operating while SnapLogic pipelines are built and validated. Each workflow transitions via a controlled switch, typically 2–4 weeks per cluster. For Choco's 1,500 workflows, a phased migration over 6–9 months is realistic.

06 — ERP Connectivity at Scale: 150+ Systems, One Platform

Choco's most significant operational challenge is the breadth of ERP systems it must connect to across its distributor base. No two distributors run the same stack. This section explains how SnapLogic addresses it.

The Choco ERP Connectivity Challenge

SCALE

150+ distinct ERP systems, each with unique API protocols, data models, and authentication methods — growing as Choco adds distributors.

FRAGMENTATION

The food distribution ERP market is highly fragmented — from global SAP deployments to food-industry-specific legacy systems (Produce Pro, CODA, Advanced Distribution) to custom-built internal platforms.

RELIABILITY

Each ERP connection is a production dependency. A failed order transmission is a failed delivery. Error handling, retry logic, and dead-letter queues are not optional features — they are table stakes.

Tier 1: Pre-built Native Snaps

For the most common ERP systems in Choco's network, SnapLogic provides certified, maintained Snaps that handle the protocol complexity, authentication, pagination, and error responses natively.

- **SAP** — BAPI/RFC, OData, IDOC, BAPIs for SD/MM/FI
- **Oracle NetSuite** — SuiteQL, REST Web Services, SOAP
- **Microsoft Dynamics 365 / BC** — OData v4 native
- **Sage 100/200/X3** — REST API and SOAP variants
- **Infor M3 / LN** — REST and SOAP
- **Acumatica / Epicor** — REST API and OData Snaps

Tier 2: Protocol-Level Snaps for Long-Tail ERPs

For the long tail of food-industry-specific and legacy ERPs that lack modern APIs, SnapLogic's protocol-native Snaps handle the transport layer — eliminating the need to write custom code for each system.

- **REST / HTTP Client** — any REST API, full OAuth 2.0 support
- **SOAP** — WSDL import and auto-generation of request structures
- **EDI** — X12 and EDIFACT parsing built in
- **FTP / SFTP / AS2** — file-based integrations natively
- **CSV / XML / JSON parsing** — format transformation native
- **Database Snaps** — direct SQL to any JDBC-compatible database

Error Handling Strategy for High-Volume ERP Integration

At 12M executions/month, failures are a statistical certainty. SnapLogic handles them via configurable retry policies per snap (fixed interval, exponential backoff, max attempts) and immediate error routing to Teams, Slack, or email with full context — which ERP, which order, which error code. No silent failures.

07 — SnapLogic vs. Alternatives — Competitive Landscape

As Choco evaluates its integration platform options, several alternatives will naturally appear in the conversation. This section provides an honest comparison framework against the most likely alternatives, focused specifically on the requirements that matter for Choco's situation.

Head-to-Head Comparison — Choco's Key Requirements

Requirement	SnapLogic	n8n (Cloud)	Make / Workato	MuleSoft	Boomi
12M+ executions/month, elastic	✓ Native elastic scaling	~ Requires custom infra	~ Depends on plan tier	✓ Enterprise grade	✓ Cloud-native elastic
150+ ERP connectors (pre-built)	✓ 800+ Snaps incl. ERP-native	× Community nodes, unmaintained	~ Good SaaS, weak ERP	✓ Extensive connector library	✓ 200+ connectors incl. ERP
EU data residency / Groundplex	✓ Groundplex on-prem/EU	~ Self-hosted only	× US-first infrastructure	✓ CloudHub EU regions	~ Atom on-prem; EU region
Native multi-tenancy	✓ Built-in org isolation	× Requires custom engineering	~ Workato yes, Make no	✓ Anypoint orgs	~ Supported via Atom config
Enterprise RBAC / SSO / SAML	✓ Full enterprise IdP support	× Cloud only; limited OSS	~ Varies by plan	✓ Full enterprise	✓ Full enterprise
AI-accelerated pipeline creation	✓ SnapGPT + Iris AI native	~ Basic AI assist (2024)	~ Emerging AI features	× Limited AI-native dev	~ Boomi GPT (less mature)
Streaming for large file processing	✓ Core architecture	× Memory-bound on large files	× Not streaming-native	✓ Streaming supported	~ Limited high-vol streaming
Total cost at Choco scale	✓ Flat-rate enterprise model	× Engineering overhead + infra	× Per-operation pricing spikes	× High licence + PS cost	× High licence + PS model

Why Not MuleSoft?

MuleSoft is technically capable of meeting Choco's requirements but comes with three structural disadvantages at Choco's scale and stage:

- **Time to value**: MuleSoft projects require specialised MuleSoft developers (Mule 4 / DataWeave). Choco's team would face a steep learning curve or need to hire/contract MuleSoft specialists — typical implementations take 6–18 months to reach production scale.
- **Cost structure**: MuleSoft's pricing model (vCores + professional services) can easily reach 3–5× SnapLogic's cost at equivalent processing volumes. For a growth-stage company, this is a significant constraint.
- **Agility**: MuleSoft's heavy governance model was designed for large enterprises with dedicated integration teams. Choco's lean engineering team needs to move faster than MuleSoft's model supports.

Why Not Make or Workato?

Both Make and Workato are strong for SaaS-to-SaaS automation but show weaknesses at Choco's specific requirements:

- **ERP connectivity**: Both platforms excel at cloud SaaS connectors (Salesforce, HubSpot, Slack). Their ERP coverage — especially for food-industry-specific ERPs and legacy systems using EDI, SOAP, or SFTP — is significantly weaker than SnapLogic's native Snap catalog.
- **Per-operation pricing**: At 12M+ executions/month, both Make's and Workato's pricing models become cost-prohibitive. Choco would face unpredictable bills as order volumes grow.
- **Streaming / large files**: Neither platform is designed for streaming large file processing — a core requirement for Choco's ERP payload handling.

Why Not Boomi?

Boomi (Dell Boomi) is a technically capable enterprise iPaaS that competes in the same tier as SnapLogic and MuleSoft. For Choco's specific situation, however, three structural factors work against it:

- **Implementation model**: Boomi projects require certified Boomi specialists. Choco's engineering team would face a steep ramp — or the cost of external consultants — before building meaningful productivity. Time-to-value is typically 4–9 months before production-scale workflows are stable.
- **Cost at Choco's scale**: Boomi's licensing model (connection-based + professional services) becomes expensive quickly at 150+ ERP connections. Like MuleSoft, the total cost of ownership at Choco's scale is materially higher than SnapLogic's flat-rate enterprise model.
- **AI maturity**: Boomi GPT exists but is less mature and less deeply integrated than SnapLogic's SnapGPT. For a team looking to accelerate new ERP integrations with AI-assisted mapping and pipeline generation, SnapLogic has a meaningful lead in production-ready AI tooling.

08 — AI-Ready Integration: SnapGPT, Agents & Agentic Automation

Integration platforms are evolving from passive data movers to active participants in intelligent workflows. SnapLogic's AI roadmap — already in production — positions Choco for the automation architecture that will differentiate food commerce platforms in 2026–2028.

🤖 SnapGPT — AI-Assisted Pipeline Development

SnapGPT is SnapLogic's native large language model integration, trained on pipeline patterns and the full Snap catalog. For Choco's engineering team, SnapGPT means:

- Describe an integration in natural language → SnapGPT generates the pipeline structure
- 65%+ reduction in time to build new ERP integration pipelines (based on SnapLogic customer benchmarks)
- Auto-mapping between source and destination schemas — critical for ERP data model translation
- Pipeline explanation and documentation generated automatically
- Error diagnosis: SnapGPT can explain why a pipeline failed and suggest fixes

For Choco's team onboarding a new distributor ERP: instead of a developer manually mapping order fields between Choco's data model and the distributor's ERP schema, SnapGPT proposes the mapping based on field names, types, and sample data — reducing a multi-day task to hours.

🔗 MCP Server — Pipelines as AI Agent Tools

SnapLogic's MCP (Model Context Protocol) Server allows AI agents and LLM-powered applications to invoke SnapLogic pipelines as tools — enabling a new class of agentic automation.

- Any SnapLogic pipeline can be exposed as an MCP tool with a natural-language description
- AI agents (Claude, GPT-4, Llama) can discover and call these tools autonomously
- Enables Choco to build AI assistants that can check order status, reroute deliveries, notify distributors — all via natural language
- Full audit trail of every agent-invoked pipeline execution

Choco use case example: A restaurant manager asks "Was my Tuesday morning order confirmed by Milano Foods?" — an AI agent queries the SnapLogic pipeline via MCP, retrieves the order confirmation from Milano Foods' ERP, and returns the answer in natural language. No engineering work required once the pipeline is built.

🚀 Agentic Integration — The 2026–2028 Horizon for Choco

Predictive Onboarding

AI agents that analyse a new distributor's ERP API documentation and automatically generate a first-draft SnapLogic integration pipeline — reducing distributor onboarding from weeks to days.

Anomaly Detection

AI monitoring of order flow patterns — detecting unusual spikes, drops, or latency increases per distributor ERP before they become SLA violations. Proactive intervention, not reactive incident response.

Natural Language Operations

Operations team members querying integration health, order status, and ERP connectivity in plain language — powered by MCP-connected pipelines feeding an LLM-based ops assistant.

09 — EU Compliance, Security & Data Residency

For a Berlin-headquartered company processing food order data across Europe, compliance is not a procurement checkbox — it is a commercial and legal obligation. This section outlines how SnapLogic addresses Choco's EU compliance requirements.

🇪🇺 GDPR & EU Data Residency

SnapLogic's Groundplex architecture enables fully EU-resident processing — data never leaves the EU unless explicitly configured to do so.

- **Groundplex**: SnapLogic's on-prem/private cloud processing node. Choco can run Groundplex on EU cloud infrastructure (AWS Frankfurt, Azure Netherlands) — all pipeline execution happens within the Groundplex, never in SnapLogic's US cloud.
- **Control plane separation**: Only metadata (execution stats, logs) flows to the SnapLogic control plane. Actual data payloads remain within the Groundplex boundary.
- **Data classification**: Pipeline parameters can enforce data classification rules — PII fields masked, restricted fields not logged.
- **DPA available**: SnapLogic provides a Data Processing Agreement (DPA) conformant with GDPR Article 28 requirements.

🔒 Security Certifications

SnapLogic maintains enterprise-grade security certifications that provide Choco with independent assurance of the platform's security posture:

- **SOC 2 Type II** — annual third-party audit of security, availability, processing integrity, confidentiality, and privacy controls
- **ISO 27001** — international information security management standard
- **GDPR-compliant** — EU-region hosting, DPA, data subject request support
- **Penetration testing** — annual third-party pen tests; results available under NDA
- **TLS 1.2+ in transit**, AES-256 at rest — all credential storage encrypted

These certifications are relevant for Choco's own enterprise sales cycles — demonstrating that the integration infrastructure underlying the Choco platform meets institutional buyer standards.

Access Control & Governance for Distributed Teams

RBAC at Every Layer

Project-level, pipeline-level, and snap-level permissions. A developer can be granted access to build pipelines in a specific ERP project without access to credentials or production execution.

SSO / SAML / OIDC

Full integration with enterprise identity providers (Okta, Azure AD, Google Workspace). Single sign-on with forced MFA. No separate SnapLogic credential management required.

Credential Vault

All ERP credentials stored encrypted in SnapLogic's credential vault — never in pipeline code. Developers reference credential aliases; only authorised users can view or modify actual credentials.

10 — 3-Phase Integration Roadmap

A phased approach reduces risk and demonstrates value quickly, starting with a Solutions Engineering-supported POV (1–2 weeks, 1–2 use cases). A successful POV leads to a buying decision and triggers the full migration phases below.

1 PHASE 1 · WEEKS 1–2

Proof of Value (POV) — Prove Additional Value

Scope

- 1–2 use cases, selected jointly during scoping workshop
- Reduced scope by design — not a production migration
- Supported by SnapLogic Solutions Engineering
- Non-production environment — validating capability, not replacing n8n

Deliverables

- Working POV pipelines in a SnapLogic-provisioned environment
- Demonstrated value across the 1–2 agreed use cases
- EU data residency architecture validated with Groundplex
- POV readout presentation with findings and recommendation
- Commercial offer for full engagement package (if POV is successful)

Success Criteria

- POV use cases demonstrably solved by SnapLogic
- Additional value vs. current n8n approach clearly evidenced
- Choco stakeholders aligned on the path forward
- Go/No-Go buying decision at end of POV

2 PHASE 2 · MONTHS 4–9

Scale — Migrate Core ERP Integrations

Scope

- Migrate top 50 highest-volume ERP integrations
- Build standardised ERP onboarding pipeline template
- Establish Choco SnapLogic Centre of Excellence (2–3 people)

Deliverables

- 50+ ERP pipelines live in production
- Standardised distributor onboarding playbook
- SnapGPT-accelerated ERP mapping library
- Onboarding time per new ERP reduced to < 3 days

Success Criteria

- 50 ERPs migrated with zero production incidents
- Distributor onboarding time halved vs n8n baseline
- Full observability dashboard operational

3 PHASE 3 · MONTHS 10–18

AI-Augmented Operations

Scope

- Complete migration of all 1,500+ workflows
- Deploy AI-assisted distributor onboarding via SnapGPT
- Build MCP-connected operations assistant
- Predictive ERP health monitoring with anomaly detection

Deliverables

- Full n8n decommission — SnapLogic sole integration platform
- AI onboarding: new ERP integration in < 1 day (AI-assisted)
- Agentic pipeline for automatic ERP error remediation

Success Criteria

- 150+ ERP connections fully migrated and stable
- Distributor onboarding fully self-serve for tier-1 ERPs
- Ready for 50M+ executions/month without re-architecture

11 — Recommended Next Steps

The following steps are designed to move from strategic alignment to a clear buying decision via a focused Proof of Value — starting with proper scoping so Choco receives the right commercial offer before any work begins.

1. **Scoping & Sizing Workshop (Week 1)**
   A structured session with Choco's engineering and commercial leads to define the 1–2 POV use cases, agree on the desired business outcome, and size the engagement correctly. This step is required before any commercial offer is made — the package must match Choco's objectives. Output: a signed-off POV scope and a SnapLogic commercial proposal for the right engagement package.

2. **Commercial Offer & POV Agreement (Week 2)**
   SnapLogic presents a formal POV package offer based on the sizing output — covering the Solutions Engineering engagement, environment provisioning, and scope of the 1–2 use cases. Choco signs off on the POV terms. This is the commercial gate before execution begins.

3. **POV Execution, Validation & Readout (Weeks 3–4)**
   SnapLogic Solutions Engineering supports the delivery of the 1–2 agreed use cases. This is a non-production evaluation — proof-of-capability, not a production migration. SnapGPT is demonstrated throughout. At the end of the POV, SnapLogic presents findings to Choco's technical and executive stakeholders. A successful POV leads directly to a buying decision and commercial negotiation for the full migration package.

WHY ACT NOW

- Choco's integration complexity grows with every new distributor onboarded — the longer the migration is deferred, the larger the migration scope becomes
- At 12M+ executions/month, every engineering hour spent managing n8n's scaling limitations is an hour not spent on product differentiation
- The AI capabilities (SnapGPT, MCP Server) that will define the next generation of distributor onboarding are available today — not in a future roadmap
- EU compliance requirements are tightening, not loosening — the Groundplex architecture provides a defensible, auditable data residency position from day one

12 — Enterprise Operations: Error Handling, Monitoring & Complex Workflow Management

Three topics that could not be fully explored during the demo — but are highly relevant to Choco's scale and complexity. This section covers how SnapLogic handles production failures, provides operational visibility, and manages large, multi-layered integration workflows at enterprise grade, drawing on the **Sigma Framework** best practices.

12.1 — Error Handling: From Failure to Remediation

At Choco's scale — 150+ ERPs, 12M+ monthly executions — production errors are not edge cases. They are a certainty. The Sigma Framework mandates a structured, reusable error handling architecture rather than ad-hoc per-pipeline logic.

Dedicated Error Pipelines

SnapLogic separates error handling from business logic. A dedicated **Error Pipeline** captures failed documents from any integration pipeline, processes them uniformly (logging, classification, alerting), and routes them for remediation — without touching the core pipeline code. This is the Sigma-recommended pattern for enterprise deployments.

Retry Policies & Dead Letter Queues

Every Snap and pipeline supports configurable retry logic: max attempts, backoff strategy (fixed or exponential), and conditional retries based on error type. Persistent failures are redirected to a **Dead Letter Queue** pattern — failed records written to durable storage (S3, Kafka, or database) for audit, replay, or manual triage. No message is silently dropped.

Choco-Specific Application

- **ERP connectivity failures**: When an SAP or Oracle ERP connection drops mid-sync, the pipeline routes the failed batch to a DLQ with full context (distributor ID, order set, timestamp), alerts the ops team, and auto-retries on the next cycle
- **Order data validation errors**: Malformed order documents from distributor ERPs are captured, logged with the source field and error reason, and held for manual review — never silently discarded
- **API gateway failures**: Rate limit or auth failures trigger exponential backoff retry; after max retries, a Slack/email alert fires to the integration ops team

SIGMA FRAMEWORK PRINCIPLE

Standardize error messages to be clear and actionable. Enforce environment separation (dev / test / prod) so error handling logic is tested before it reaches production. Maintain error pipelines in version control (Git integration) to enable rollback if error handling logic itself needs fixing.

12.2 — Advanced Monitoring: Full Observability at Production Scale

Monitoring is not just about knowing when something broke — it is about knowing before it breaks, and understanding why when it does. SnapLogic provides multiple layers of observability, from built-in dashboards to full enterprise APM integration.

Built-in SnapLogic Dashboard

- Real-time pipeline execution status across all projects and environments
- Error rate trends, execution time percentiles, and throughput metrics per pipeline
- Per-Snap performance breakdown — identify bottlenecks without guessing
- Historical execution logs with full document-level traceability
- Alert configuration: threshold-based notifications via email or webhook

OpenTelemetry + Datadog Integration

- SnapLogic supports **OpenTelemetry** export — pipeline telemetry (traces, metrics, logs) streamed to any OTel-compatible backend
- Datadog is the primary enterprise target: custom dashboards correlating SnapLogic execution data with infrastructure and application metrics
- Enables **cross-system correlation** — link a pipeline failure to an upstream ERP API degradation in the same Datadog dashboard
- Historical data retained for compliance, SLA reporting, and post-incident analysis

Choco Monitoring Blueprint

- **Critical KPIs to track**: pipeline failure rate per ERP, execution latency p95, order sync lag, error recurrence rate
- **Alerting tiers**: P1 (ERP sync down >5 min → PagerDuty), P2 (error rate spike → Slack ops channel), P3 (degraded throughput → daily digest)
- **Capacity signals**: Groundplex worker CPU/memory surfaced in Datadog; auto-scale triggers defined ahead of growth spikes
- **SLA dashboards**: per-distributor uptime and order processing compliance, exportable for partner reporting

12.3 — Managing Complex Workflows: A Complete Example

The UI demo surfaced a valid question: how does SnapLogic actually manage a complex, multi-step integration without turning into an unmaintainable tangle? The answer is **modular pipeline architecture** — the cornerstone of the Sigma Framework for enterprise integration.

The Modular Pattern (Sigma Framework)

SnapLogic pipelines are designed to be **nested, not monolithic**. Complex workflows are broken into reusable child pipelines called via the **Pipeline Execute Snap**. Each child pipeline owns one concern: authentication, data transformation, error handling, or notification. The parent pipeline is a clean orchestration layer — readable, testable, maintainable.

- **Child pipelines**: reusable across workflows; tested independently
- **Expression libraries**: shared transformation logic, version-controlled
- **Environment variables**: no hardcoded endpoints — promotion from dev → test → prod without code changes
- **Shared component library**: governed by CoE; any team can consume, no team can break it silently

WHY THIS SCALES TO 150+ ERPS

Each ERP type (SAP B1, SAP ECC, Oracle, Microsoft BC, etc.) has its own Auth and Extract child pipeline. But the Transform, Validate, Load, and Notify pipelines are **shared across all ERP types**. Adding a new ERP connector means writing one Auth + one Extract pipeline — not rebuilding the entire workflow. This is the architecture that takes onboarding from weeks to days.

⚙️ **Governance via Center of Excellence (CoE):** The Sigma Framework recommends a CoE as the governing body for shared pipeline components. The CoE owns the canonical pipelines (Transform, Validate, Load) and defines the contribution model for new ERP-specific connectors. This prevents duplication, enforces standards, and ensures production quality across all 150+ integration workflows — regardless of which team built the ERP-specific pieces.

Example: Distributor ERP Onboarding Workflow

A new distributor connects their SAP B1 instance to Choco's platform. This single business event triggers a multi-step workflow — here is how it is structured in SnapLogic:

Layer	Child Pipeline	Responsibility
Orchestrator	erp-onboard-main	Sequences all steps; owns retry and rollback logic
Auth	sap-b1-auth	Handles credential retrieval from vault + session token; reused by 80+ other SAP pipelines
Extract	erp-schema-discovery	Pulls distributor ERP schema; validates required fields
Transform	order-canonical-map	Maps ERP-specific order format to Choco canonical model; shared pipeline used across all distributors
Validate	data-quality-check	Business rules validation; routes failures to error pipeline automatically
Load	choco-platform-write	Writes to Choco platform API; idempotent — safe to retry without duplicate records
Notify	onboard-notification	Sends confirmation to ops team + distributor; logs to audit trail

OPERATIONAL READINESS SUMMARY FOR CHOCO

- **Error handling** — Dedicated error pipelines, retry with backoff, DLQ for all failed records. No silent failures. Full auditability.
- **Monitoring** — SnapLogic native dashboard + OpenTelemetry export to Datadog. Alerts tiered by severity. SLA tracking per distributor.
- **Complex workflows** — Modular pipeline architecture. Shared canonical pipelines reused across all ERPs. CoE governance for shared components. New ERP onboarding in days, not weeks.
- **All three capabilities** are production-proven in comparable enterprise deployments and are available from day one of the SnapLogic engagement.